

QINGKAI SHI

Email: qingkaishi@gmail.com **Homepage:** <https://qingkaishi.github.io/> **Twitter:** @QingkaiS

EDUCATION AND EMPLOYMENT

- **Purdue University** 2021 - Now
Postdoc Research Associate
- **Sourcebrella Inc.** 2015 - 2021
Co-founder
Sourcebrella Inc. commercialized my research on static bug finding
Sourcebrella Inc. was acquired by Ant Group in 2020
- **The Hong Kong University of Science and Technology** 2015 - 2020
Ph.D. in Computer Science and Engineering
- **Nanjing University & MoocTest Inc.** 2012 - 2015
Research Engineer
- **Nanjing University** 2008 - 2012
B.S. in Software Engineering

RESEARCH HIGHLIGHTS

- My research interest centers around programming language, software engineering, and cybersecurity, focusing on the use of compiler techniques, including both static and dynamic program analysis, for ensuring software security and reliability.
- My research has been commercialized and deployed in many Global 500 companies, helping them hunt deeply-hidden software vulnerabilities. The startup, Sourcebrella Inc., that commercializes my research has been acquired by Ant Group.
- My research has won me an ACM SIGSOFT Distinguished Paper Award (2019), three champions in NASAC prototype competitions (2016, 2018a, 2018b), Hong Kong PhD Fellowship (2015), and China National Scholarship (2010, 2015).

PUBLICATIONS

- [1] Heqing Huang, Yiyuan Guo, **Qingkai Shi***, Peisen Yao, Rongxin Wu, and Charles Zhang. Beacon: Directed Grey-Box Fuzzing with Provable Path Pruning. In *Proceedings of the 43rd IEEE Symposium on Security and Privacy (S&P'22)*. IEEE, 2022. (* corresponding author and papers under my supervision)
- [2] **Qingkai Shi**, Yongchao Wang, and Charles Zhang. Indexing Context-Sensitive Reachability. In *Proceedings of the 36th ACM SIGPLAN Conference on Object Oriented Programming, Systems, Languages and Applications (OOPSLA'21)*. ACM, 2021.
- [3] Peisen Yao, **Qingkai Shi***, Heqing Huang, and Charles Zhang. Program Analysis via Efficient Symbolic Abstraction. In *Proceedings of the 36th ACM SIGPLAN Conference on Object Oriented Programming, Systems, Languages and Applications (OOPSLA'21)*. ACM, 2021.

- [4] Peisen Yao, Heqing Huang, Wensheng Tang, **Qingkai Shi**, Rongxin Wu, and Charles Zhang. Skeletal Approximation Enumeration for SMT Solver Testing. In *Proceedings of the 29th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE'21)*. pp 1141–1153. ACM, 2021. <https://doi.org/10.1145/3468264.3468540>
- [5] Peisen Yao, Heqing Huang, Wensheng Tang, **Qingkai Shi**, Rongxin Wu, and Charles Zhang. Fuzzing SMT Solvers via Two-Dimensional Input Space Exploration. In *Proceedings of the 30th ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA'21)*. pp 322–335. ACM, 2021. <https://doi.org/10.1145/3460319.3464803>
- [6] **Qingkai Shi**, Peisen Yao, Rongxin Wu, and Charles Zhang. Path-Sensitive Sparse Analysis without Path Conditions. In *Proceedings of the 42nd ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI'21)*. pp 930–943. ACM, 2021. <https://doi.org/10.1145/3453483.3454086>
- [7] Peisen Yao, **Qingkai Shi***, Heqing Huang, and Charles Zhang. Fast Bit-Vector Satisfiability. In *Proceedings of the 29th ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA'20)*. pp 38–50. ACM, 2020. <https://doi.org/10.1145/3395363.3397378>
- [8] Gang Fan, Chengpeng Wang, Rongxin Wu, Xiao Xiao, **Qingkai Shi**, and Charles Zhang. Escaping Dependency Hell: Finding Build Dependency Errors with the Unified Dependency Graph. In *Proceedings of the 29th ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA'20)*. pp 463–474. ACM, 2020. <https://doi.org/10.1145/3395363.3397388>
- [9] Chunrong Fang, Zixi Liu, Yangyang Shi, Jeff Huang, and **Qingkai Shi**. Functional Code Clone Detection with Syntax and Semantics Fusion Learning. In *Proceedings of the 29th ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA'20)*. pp 516–527. ACM, 2020. <https://doi.org/10.1145/3395363.3397362>
- [10] Yang Feng, **Qingkai Shi***, Xinyu Gao, Jun Wan, Chunrong Fang, and Zhenyu Chen. DeepGini: Prioritizing Massive Tests to Enhance the Robustness of Deep Neural Networks. In *Proceedings of the 29th ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA'20)*. pp 177–188. ACM, 2020. <https://doi.org/10.1145/3395363.3397357>
- [11] **Qingkai Shi** and Charles Zhang. Pipelining Bottom-up Data Flow Analysis. In *Proceedings of the 42nd ACM/IEEE International Conference on Software Engineering (ICSE'20)*. pp 835–847. ACM, 2020. <https://doi.org/10.1145/3377811.3380425>
- [12] **Qingkai Shi**, Rongxin Wu, Gang Fan, and Charles Zhang. Conquering the Extensional Scalability Problem for Value-Flow Analysis Frameworks. In *Proceedings of the 42nd ACM/IEEE International Conference on Software Engineering (ICSE'20)*. pp 812–823. ACM, 2020. <https://doi.org/10.1145/3377811.3380346>
- [13] Heqing Huang, Peisen Yao, Rongxin Wu, **Qingkai Shi***, and Charles Zhang. Pangolin: Incremental Hybrid Fuzzing with Polyhedral Path Abstraction. In *Proceedings of the 41st IEEE Symposium on Security and Privacy (S&P'20)*. pp 1613–1627. IEEE, 2020. <https://doi.org/10.1109/SP40000.2020.00063>
- [14] Gang Fan, Rongxin Wu, **Qingkai Shi**, Xiao Xiao, Jinguo Zhou, and Charles Zhang. SMOKE: Scalable Path-Sensitive Memory Leak Detection for Millions of Lines of Code. In *Proceedings of the 41st ACM/IEEE International Conference on Software Engineering (ICSE'19)*. pp 72–82. IEEE, 2019. <https://doi.org/10.1109/ICSE.2019.00025> **ACM SIGSOFT Distinguished Paper Award**
- [15] **Qingkai Shi**, Xiao Xiao, Rongxin Wu, Jinguo Zhou, Gang Fan, and Charles Zhang. Pinpoint: Fast and Precise Sparse Value Flow Analysis for Million Lines of Code. In *Proceedings of the 39th ACM SIGPLAN Conference*

on *Programming Language Design and Implementation (PLDI'18)*. pp 693–706. ACM, 2018. <https://doi.org/10.1145/3192366.3192418>

- [16] **Qingkai Shi**, Zhenyu Chen, Chunrong Fang, Yang Feng, and Baowen Xu. Measuring the Diversity of a Test Set with Distance Entropy. In *IEEE Transactions on Reliability (TRel'16)*, Vol. 65, No. 1. pp 19–27. IEEE, 2016. <https://doi.org/10.1109/TR.2015.2434953>
- [17] **Qingkai Shi**, Jeff Huang, Zhenyu Chen, and Baowen Xu. Verifying Synchronization for Atomicity Violation Fixing. In *IEEE Transactions on Software Engineering (TSE'16)*, Vol. 42, No. 3. pp 280–296. IEEE, 2016. <https://doi.org/10.1109/TSE.2015.2477820>

INVITED TALKS AND PRESENTATIONS

- “Path-Sensitive Sparse Analysis without Path Conditions”, the 42nd ACM SIGPLAN conference on Programming Language Design and Implementation (PLDI'21), Virtual, June, 2021. **Presentation.**
- “Taming Path-Sensitivity for Static Code Analysis on Industrial Scale”, Nanjing University, Nanjing, China, April, 2021. **Invited Talk.**
- “Pipelining Bottom-up Data Flow Analysis”, the 42nd ACM/IEEE International Conference on Software Engineering (ICSE'20), Virtual, July, 2020. **Presentation.**
- “Conquering the Extensional Scalability Problem for Value-Flow Analysis Frameworks”, the 42nd ACM/IEEE International Conference on Software Engineering (ICSE'20), Virtual, July, 2020. **Presentation.**
- “The Next Generation Software Security Analysis”, Nanjing University, Nanjing, China, May, 2020. **Invited Talk.**
- “Marrying Taint Analysis with Big Data Analytics”, Southern University of Science and Technology, Shenzhen, China, November, 2018. **Invited Talk.**
- “Fast and Precise Sparse Value Flow Analysis for Million Lines of Code”, the 39th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI'18), Philadelphia, the United States, June, 2018. **Presentation.**

CVE IDENTIFIERS

- [1] CVE-2017-14739: ImageMagick 7.0.7–4 mishandles failed memory allocation, which allows remote attackers to cause a denial of service.
- [2] CVE-2017-14952: International Components for Unicode (ICU) for C/C++ through 59.1 contains a double free that allows remote attackers to execute arbitrary code.
- [3] CVE-2017-15096: GlusterFS in versions prior to 3.10 contains a null pointer dereference that may cause denial of service.
- [4] CVE-2017-16892: Bftpd 4.6 contains a memory leak which occurs if a mal-crafted sequence of FTP requests are received.
- [5] CVE-2017-1000445: ImageMagick 7.0.7–1 and older version are vulnerable to null pointer dereference in the MagickCore component and might lead to denial of service.

- [6] CVE-2018-20786: libvterm through 0+bzr726, as used in Vim and other products, mishandles certain out-of-memory conditions, leading to a denial of service (application crash), related to screen.c, etc.
- [7] CVE-2019-13238: An issue was discovered in Bento4 1.5.1.0. A memory allocation failure is unhandled in Core/Ap4SdpAtom.cpp and leads to crashes. When parsing input video, the program allocates a new buffer to parse an atom in the stream. The unhandled memory allocation failure causes a direct copy to a NULL pointer.
- [8] CVE-2019-13959: In Bento4 1.5.1-627, AP4_DataBuffer::SetDataSize does not handle reallocation failures, leading to a memory copy into a NULL pointer.
- [9] CVE-2019-13960: In libjpeg-turbo 2.0.2, a large amount of memory can be used during processing of an invalid progressive JPEG image containing incorrect width and height values in the image header.
- [10] CVE-2020-19715: An integer overflow vulnerability in the getUShort function of Exiv2 0.27.1 results in segmentation faults within the application, leading to a denial of service (DOS).
- [11] CVE-2020-19716: A buffer overflow vulnerability in the Databuf function in types.cpp of Exiv2 v0.27.1 leads to a denial of service (DOS).
- [12] CVE-2020-19717: An unhandled memory allocation failure in Core/Ap48bdlAtom.cpp of Bento 1.5.1-628 causes a NULL pointer dereference, leading to a denial of service (DOS).
- [13] CVE-2020-19718: An unhandled memory allocation failure in Core/Ap4Atom.cpp of Bento 1.5.1-628 causes a NULL pointer dereference, leading to a denial of service (DOS).
- [14] CVE-2020-19719: A buffer overflow vulnerability in Ap4ElstAtom.cpp of Bento 1.5.1-628 leads to a denial of service (DOS).
- [15] CVE-2020-19720: An unhandled memory allocation failure in Core/AP4IkmsAtom.cpp of Bento 1.5.1-628 causes a NULL pointer dereference, leading to a denial of service (DOS).
- [16] CVE-2020-19721: A heap buffer overflow vulnerability in Ap4TrunAtom.cpp of Bento 1.5.1-628 may lead to an out-of-bounds write while running mp42aac, leading to system crashes and a denial of service (DOS).
- [17] CVE-2020-19722: An unhandled memory allocation failure in Core/Ap4Atom.cpp of Bento 1.5.1-628 causes a direct copy to NULL pointer dereference, leading to a denial of service (DOS).

PATENTS

- [1] Defect detection method, device, system, and computer readable medium. (US Patent No. 20190108003, China Patent No. 201811013103.6).
- [2] Use-after-free detection method, device, system, and computer readable medium. (China Patent No. 201811013000.X).
- [3] SQL-injection detection method, device, system, and computer readable medium. (China Patent No. 201811015751.5).
- [4] A method of obtaining the feedback on the teaching of Java unit testing. (China Patent No. 2016102941812).
- [5] Inter-procedural null dereference detection method, device, system, and computer readable medium. (China Patent No. 2018110146864).
- [6] A method of verifying synchronization for atomicity violation fixing. (China Patent No. 2014107099836).

[7] A method of randomly selecting co-diversified test cases. (China Patent No. 2012100526910).

[8] A method of controlling computer with mobile phone's inner sensors. (China Patent No. 2011104124584).

PROFESSIONAL SERVICES

- Artifact Evaluation Committee:
 - The 49th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'22), 16 - 22 January 2022, Philadelphia, Pennsylvania, the United States.
- Student Volunteer:
 - The 37th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI'16), 13 - 17 June 2016, Santa Barbara, California, the United States.
 - The 13th International Conference on Quality Software (QSIC'13), 29 - 30 July 2013, Nanjing, China.
 - The 33rd ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI'12), 11 - 16 June 2012, Beijing, China.

TEACHING EXPERIENCE

- Teaching Assistant for COMP3111/3111H: Software Engineering (Fall 2018)
- Teaching Assistant for COMP4111: Software Engineering Practices (Spring 2018)
- Teaching Assistant for COMP4111: Software Engineering Practices (Fall 2016)